



Day 6: Data Splitting and Feature Engineering

Goal

By the end of Day 7, interns will:

- Understand **why** and **how** to split data for model training and evaluation.
 - Apply **feature engineering techniques** to improve model performance.
 - Learn **how to prepare raw data** into usable inputs for machine learning models.
-

1. What is Data Splitting?

Why We Split the Dataset

In Machine Learning, we **never train and test on the same data**. Why?

- To **prevent overfitting** (learning too well on training data)
- To evaluate how well your model generalizes to new data

Types of Splits

- **Training Set** – used to train the model (usually 70–80%)
- **Test Set** – used to evaluate the model (usually 20–30%)

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

2. What is Feature Engineering?

Feature engineering is the process of **transforming raw data** into features that are more informative for your machine learning model.

3. Types of Feature Engineering

A. Feature Scaling (Standardization)

Why?

- Some algorithms (like Linear Regression, KNN) are sensitive to feature scale.

Standardization Formula:

$$z = \frac{(x - \mu)}{\sigma}$$

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

B. Polynomial Features (Interaction Terms)

Use when:

- Relationships between variables may not be linear
- Polynomial regression helps model non-linearity

```
from sklearn.preprocessing import PolynomialFeatures
```

```
poly = PolynomialFeatures(degree=2, include_bias=False)
X_poly = poly.fit_transform(X_scaled)
```

Adds combinations like:

$x_1, x_2, x_1^2, x_2^2, x_1 \cdot x_2$

C. One-Hot Encoding (for categorical variables)

Not used in Day 7 since California dataset is numeric, but here's how it works:

```
pd.get_dummies(df['Category'])
```

D. Feature Selection

Why?

- Reduce dimensionality
- Remove irrelevant or noisy features

```
from sklearn.feature_selection import SelectKBest, f_regression
```

```
selector = SelectKBest(score_func=f_regression, k=5)
```

```
X_selected = selector.fit_transform(X_scaled, y)
```

Mini Project Summary

Using the **California Housing** dataset, interns will:

1. Load the dataset
 2. Split it into train/test
 3. Standardize the features
 4. Create polynomial features
 5. Optionally select the top 5 features
-

Gamification / Intern Challenge Ideas

- **Experiment** with different degrees in PolynomialFeatures.
 - **Compare models:** Try building a Linear Regression on scaled vs. unscaled vs. poly-transformed features.
 - **Leaderboard:** Who gets the lowest MSE on test set with top 5 features?
-

Key Takeaways

- Always split data to avoid overfitting.
 - Scaling and feature engineering are critical for many ML algorithms.
 - Polynomial features help capture non-linear patterns.
 - Feature selection can boost speed and reduce overfitting.
-



Classification – Logistic Regression

Objective

Understand the fundamentals of classification and implement **logistic regression** using a real-world binary dataset.

What is Classification?

Classification is a type of **Supervised Learning** where:

- The **output** is **categorical** (e.g., Yes/No, Spam/Not Spam, Malignant/Benign)
- The goal is to **predict the correct class** for new data

Examples of Classification:

Problem	Input Features	Output
Email Spam Detection	Email content, sender	Spam / Not Spam
Tumor Classification	Cell size, nucleus shape	Malignant / Benign

Logistic Regression

Despite the name, **Logistic Regression** is a **classification algorithm**, not a regression algorithm.

Core Idea:

- Instead of predicting a continuous value, we predict a **probability between 0 and 1**
- If the predicted probability > 0.5, class = 1; else, class = 0

It uses the Sigmoid Function:

$$\sigma(z) = \frac{1}{1 + e^{-z}} \quad \sigma(z) = \frac{1}{1 + e^{-z}}$$

Where z is a linear combination of input features ($z = w \cdot x + b$)

Dataset Used: Breast Cancer Dataset

- Built-in dataset in `sklearn`
- Predict whether a tumor is **malignant** or **benign**
- Input features: Mean radius, texture, area, etc.
- Target: 0 = malignant, 1 = benign

```
from sklearn.datasets import load_breast_cancer
data = load_breast_cancer()
```

Step-by-Step Implementation

Step 1: Split the Dataset

Use `train_test_split()` to divide data into training and test sets.

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

- **Training set:** Used to train the model
 - **Test set:** Used to evaluate performance
-

Step 2: Feature Scaling

Why scale?

- Logistic regression is sensitive to feature magnitudes

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

Step 3: Train Logistic Regression Model

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train_scaled, y_train)
```

- Fits the logistic function to the training data
 - Learns the weights **w** and bias **b**
-



Step 4: Evaluate the Model

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
y_pred = model.predict(X_test_scaled)
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))  
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))  
print("Classification Report:\n", classification_report(y_test, y_pred))
```



Evaluation Metrics:

Metric	What It Measures
Accuracy	Overall correctness
Precision	How many predicted positives are correct
Recall	How many actual positives were correctly predicted
F1-Score	Balance between precision and recall



Summary

Concept	Description
Classification	Predict categories
Logistic Regression	Maps inputs to probabilities
Sigmoid Function	Converts output to 0–1
Feature Scaling	Normalizes input features
Evaluation	Accuracy, precision, recall, F1



Intern Exercise

1. Change test size to 30%. Does accuracy change?
2. Print model coefficients using `model.coef_`. Which features are most impactful?

3. Add a challenge: Try with another binary dataset like `load_digits()` (2-class version).
